

ECE 09495/09595

Advanced Emerging Topics in Computational Intelligence, Machine Learning and Data Mining
Emerging Topics in Computational Intelligence, Machine Learning and Data Mining

Ghulam Rasool, PhD

30 January 2019

Lecture 3 and 4

Housekeeping

- Google ML Crash Course
- Projects
- Google Colab

Data Preparation

May take up to 70-80% of the teams' effort initially

Data Preparation

Intelligence

Interpretation

Deployment



- Data sources
- Data labelling
- Clean up
- Visualize

- Analytics
- Machine Learning
- Deep Learning

Data Preparation

Evaluation and design iteration



Data Preparation

Intelligence

Interpretation

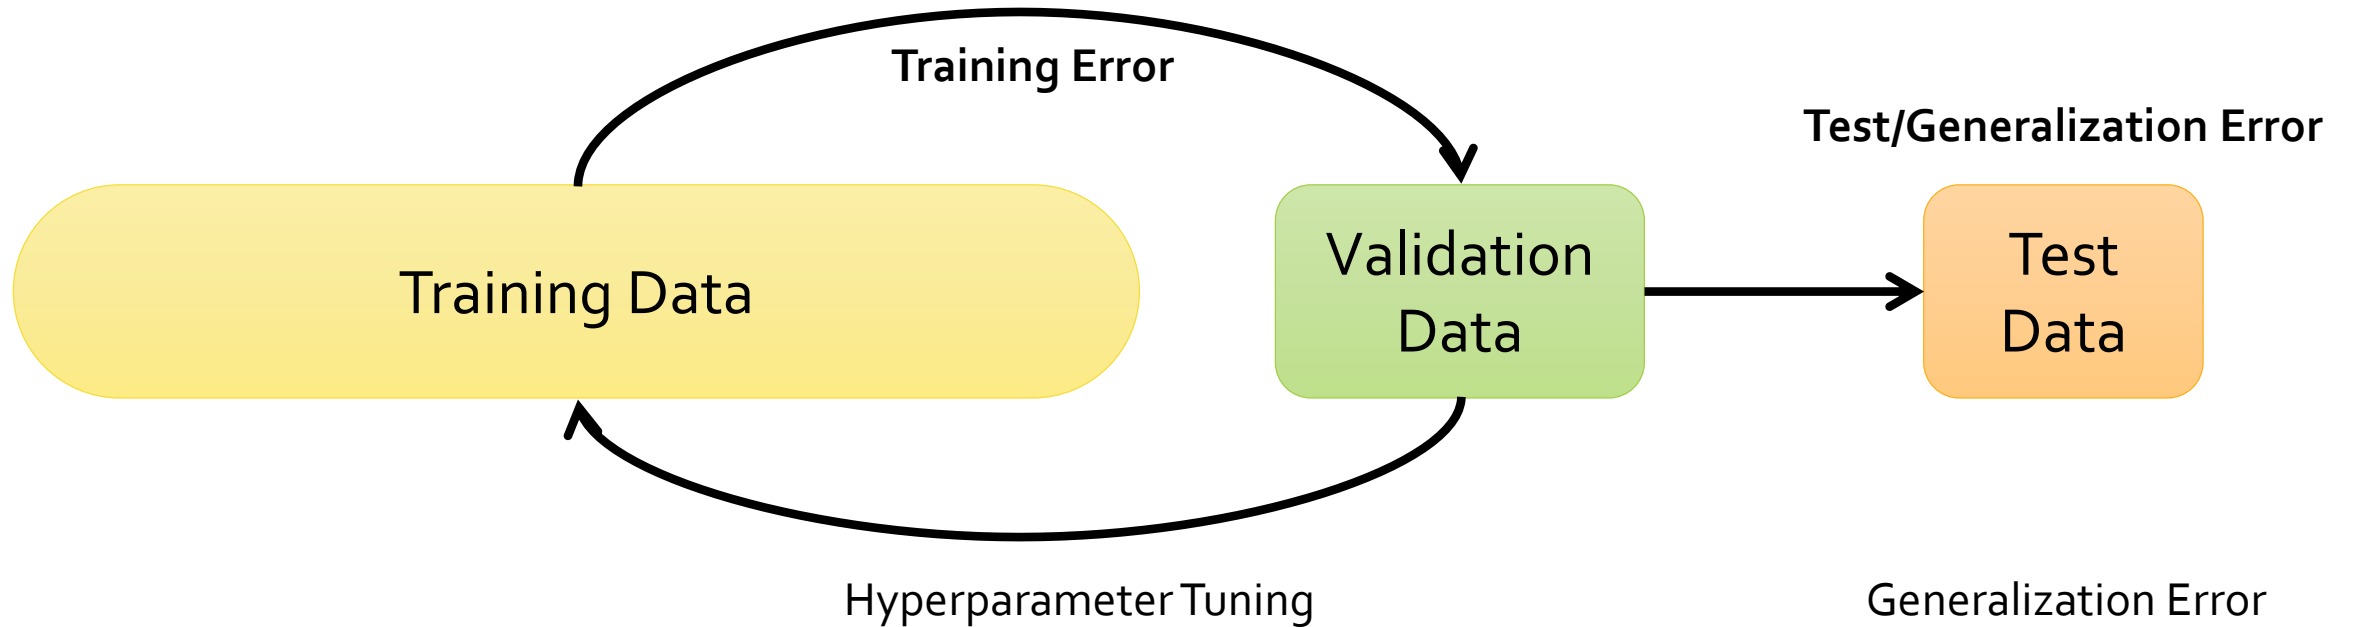
Deployment



- Data sources
- Data labelling
- Clean up
- Visualize

- Analytics
- Machine Learning
- Deep Learning

Data Distribution




- Make the training error small ->
- Make the gap between training and test error small ->

Model performs well
Model generalizes well



Underfitting

Underfitting occurs when the model is not able to obtain a sufficiently low error value on the **training set**.



Underfitting and Overfitting

Underfitting occurs when the model is not able to obtain a sufficiently low error value on the **training set**.



Overfitting occurs when the gap between the **training accuracy** and **test accuracy** is too large.

Regularization – Dropout, L1, L2, Weight Decay

Linear Models - Why?

Supervised Learning

Linear Regression

Logistic Regression

Neuron

- Labels are real numbers
- Interpretability
- Closed-form solution
- Gradient descent

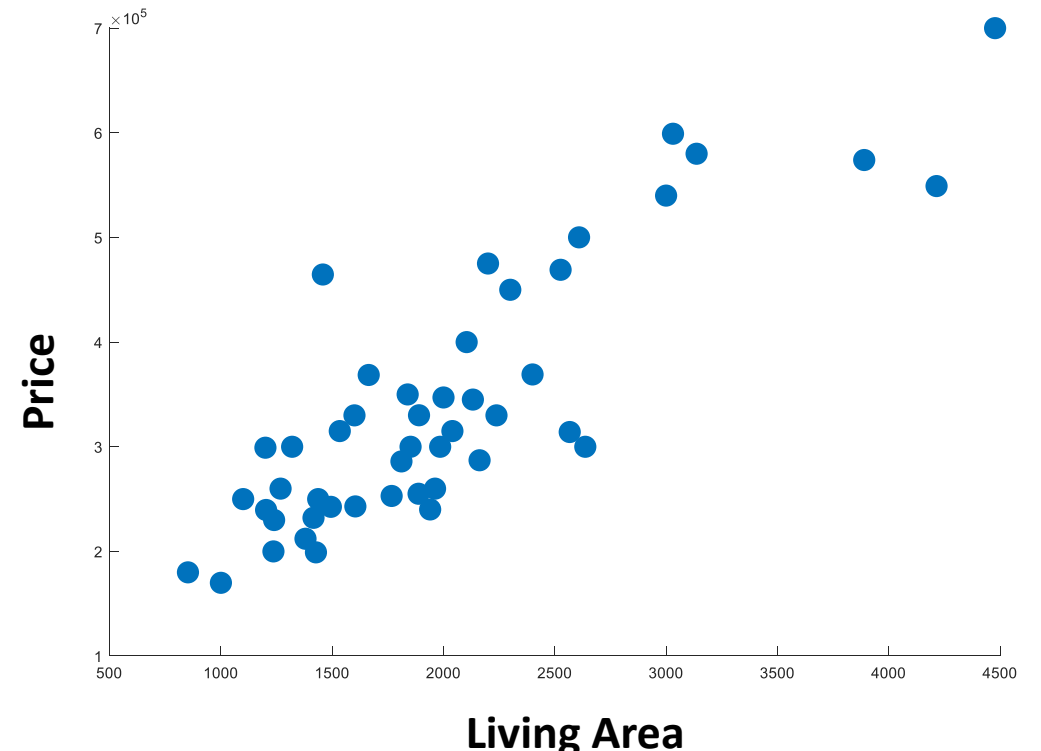
- Labels are categories
- Classification problem
- Nonlinearity
- Gradient descent

Neural networks are stacked
layers of neurons

Linear Regression

- We have a dataset giving the living areas and prices of houses

| Living Area | Price \$K |
|-------------|-----------|
| 2104 | 400 |
| 1600 | 330 |
| 2400 | 369 |
| ⋮ | ⋮ |




Linear Model

- Lets try to make a linear model

$$h_w(x) = w_0 + w_1x$$

$$y = mx + c$$



Notations

- $x^{(i)}$ is the i -th input example (feature) with $y^{(i)}$ as the target or the label.
- Together $(x^{(i)}, y^{(i)})$ are referred to as the i -th training example.
- $\{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$ is the training data set.
- \mathcal{X} is the space of input values, and \mathcal{Y} the space of labels/targets, $\mathcal{X} = \mathcal{Y} = \mathbb{R}$

Our goal

$$h : \mathcal{X} \mapsto \mathcal{Y}$$

so that $h(x)$ is a good predictor of y

More features

- We may have more features, living area and number of bedrooms, so we have $x_1^{(i)}$, and $x_2^{(i)}$. Feature space is now two-dimensional, $x \in \mathbb{R}^2$.
- Who decides how many features we have?

Lets try to make a linear model

$$h_w(x) = w_0 + w_1x_1 + w_2x_2$$

$$h(x) = \sum_{i=0}^n w_i x_i = w^T x$$

x_0 is set to 1.

w_i are the parameters or weights and w is the vector of parameters/weights. x is the vector on inputs.

What are known and what are unknowns?

The unknown w and the Least-Squares

Given training set $\{(x^{(i)}, y^{(i)}), i = 1, \dots, m\}$, how do we learn the parameters w such that $h(x)$ is close to the ground truth y .

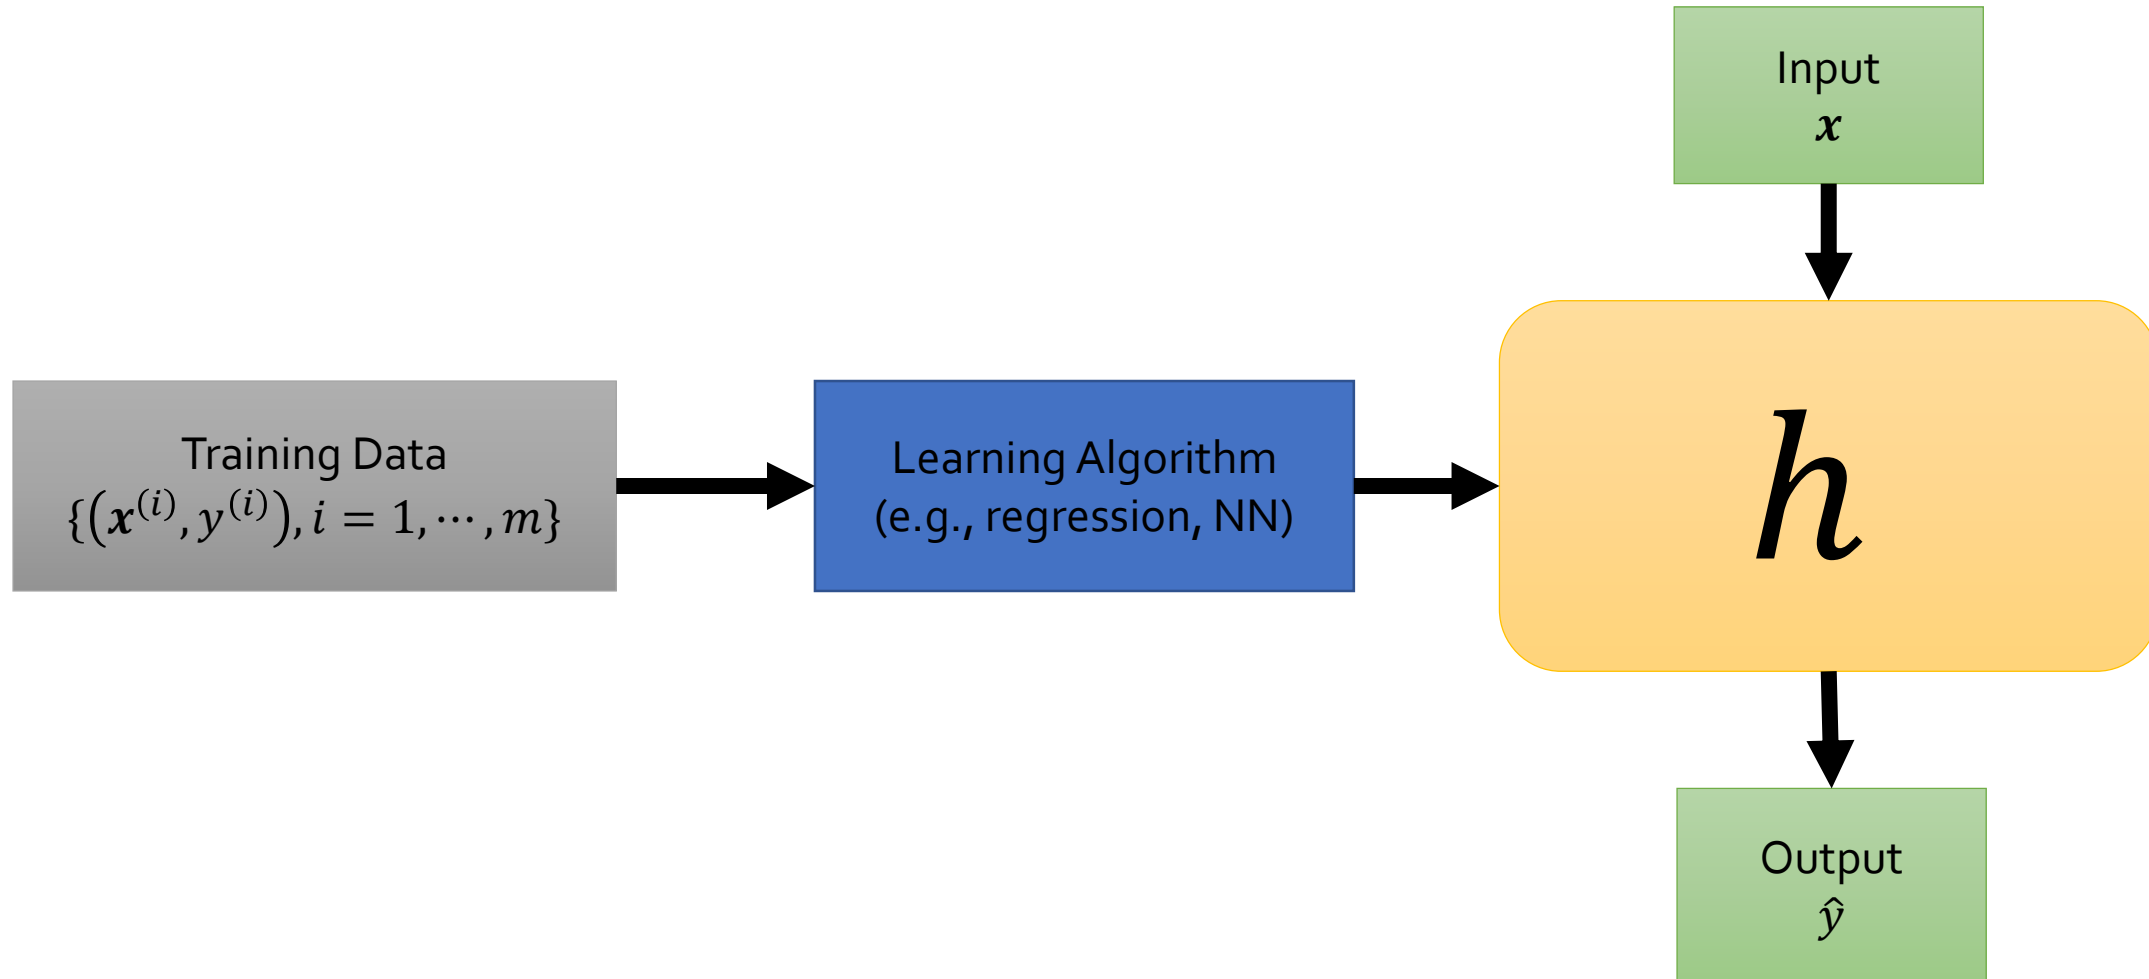
Choose some w randomly and test all examples and compare $h(x)$ and y .

Lets define cost function

$$J(w) = \frac{1}{2} \sum_{i=1}^m (h(x^{(i)}) - y^i)^2$$

This is a least-squares cost function and referred to as the ordinary least-squares (OLS) regression.

Machine Learning - Supervised Learning



Linear Regression - cost function

$$J(w) = \frac{1}{2} \sum_{i=1}^m \left(h(x^{(i)}) - y^i \right)^2$$

$$h(x) = \sum_{i=0}^n w_i x_i = w^T x$$

Square Error

Ground truth

Sum over all given examples

Goal : Minimize $J(w)$

Gradient Descent

We want to choose w so as to minimize $J(w)$

Starts with some “guess” for w , and then repeatedly change w to make $J(w)$ smaller.

The gradient descent algorithm starts with some initial w , and repeatedly performs the update:

$$w_j := w_j - \alpha \frac{\partial}{\partial w_j} J(w) \quad j = 1, \dots, n$$

α is the learning rate (more on this later)

For a single training example, we may solve and get:

$$w_j := w_j - \alpha (h_w(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

LMS (Least Mean Squares) Update Rule




Gradient descent: But how?

$$f(w) = w^2 + 20w$$

$$\frac{df(w)}{dw} = 2w + 20$$

$$\alpha = 0.1$$

$$w := w + \alpha \left(-\frac{df(w)}{dw} \right)$$



Batch gradient descent

Loop till convergence:

$$w_j := w_j - \alpha \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$\frac{\partial}{\partial w_j} J(w)$$

This method looks at every example in the entire training set on every step

Stochastic gradient descent

Loop:{

for $i = 1:m$ {

$w_j := w_j - \alpha(h_w(x^{(i)}) - y^{(i)})x_j^{(i)}$ for all j

}

}

Batch gradient descent has to scan through the entire training set before taking a single step of updating w – a costly operation if m is large

Stochastic gradient descent can start making progress right away, and continues to make progress with each example it looks at.

Stochastic gradient descent may get w close to the minimum much faster than batch gradient descent. It may never converge to the minimum, that is, w may keep oscillating around the minimum of $J(w)$

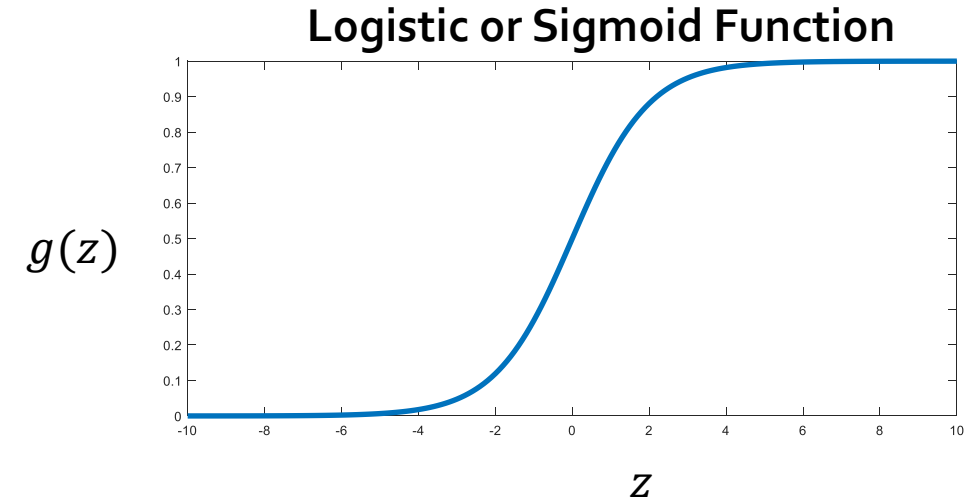

Logistic Regression

Labels are now discrete categories, e.g.
 $y \in \{0, 1\}$ – a binary classification
problem

Previously : $h_w(x) = w^T x$

Now: $h_w(x) = g(w^T x)$ such that
 $h_w(x) \in \{0, 1\}$

$$g(z) = \frac{1}{1 + \exp(-z)}$$



Logistic regression: How to estimate w ?

Lets start with probability assignment:

$$P(y = 1|x; w) = h_w(x) \text{ and thus } P(y = 0|x; w) = 1 - h_w(x)$$
$$p(y|x; w) = [h_w(x)]^y [1 - h_w(x)]^{1-y}$$

For m independent training examples:

$$L(w) = \prod_{i=1}^m [h_w(x^{(i)})]^{y^{(i)}} [1 - h_w(x^{(i)})]^{1-y^{(i)}}$$
$$l(w) = \log L(w) = \sum_{i=1}^m y^{(i)} \log h(x^{(i)}) + (1 - y^{(i)}) \log(1 - h(x^{(i)}))$$

How do we maximize it? Gradient Ascent

Logistic Regression Loss Function

Inputs: x, y

Parameters: w and b

Output : \hat{y}

$$J(w, b) = - \sum_{i=1}^m y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

$$\hat{y}^{(i)} = \sigma(w^T x^{(i)} + b) = \frac{1}{1 + \exp(-w^T x^{(i)} + b)}$$

w and b can be combined

$$w^T x + b = [w_1, w_2] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + b$$

$$= [b, w_1, w_2] \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}$$

$$y^{(i)} = 1 \text{ and } h(x^{(i)}) = 0.999 \quad \text{Loss} \approx 0$$

$$y^{(i)} = 0 \text{ and } h(x^{(i)}) = 0.001 \quad \text{Loss} \approx 0$$

$$y^{(i)} = 1 \text{ and } h(x^{(i)}) = 0.001 \quad \text{Loss} \approx 3$$

$$y^{(i)} = 0 \text{ and } h(x^{(i)}) = 0.999 \quad \text{Loss} \approx 3$$

Loss function / Cost function

- Function of the unknown weights, w , and b
- Differentiable
- Must output a scalar

Linear Regression and Logistic Regression


When the target variable (y) we are trying to predict is continuous, such as in our housing example, we call the learning problem a regression problem. We can solve this problem using linear models and this is referred to as **Linear Regression**.

Examples???



When y can take on only a small number of discrete values or categories, (e.g., low-range, mid-range, expensive, crazy-expensive), we call it a classification problem. We can solve this problem using linear models, this is referred to as **Logistic Regression**.

Examples???

Logistic Regression – 1 example case



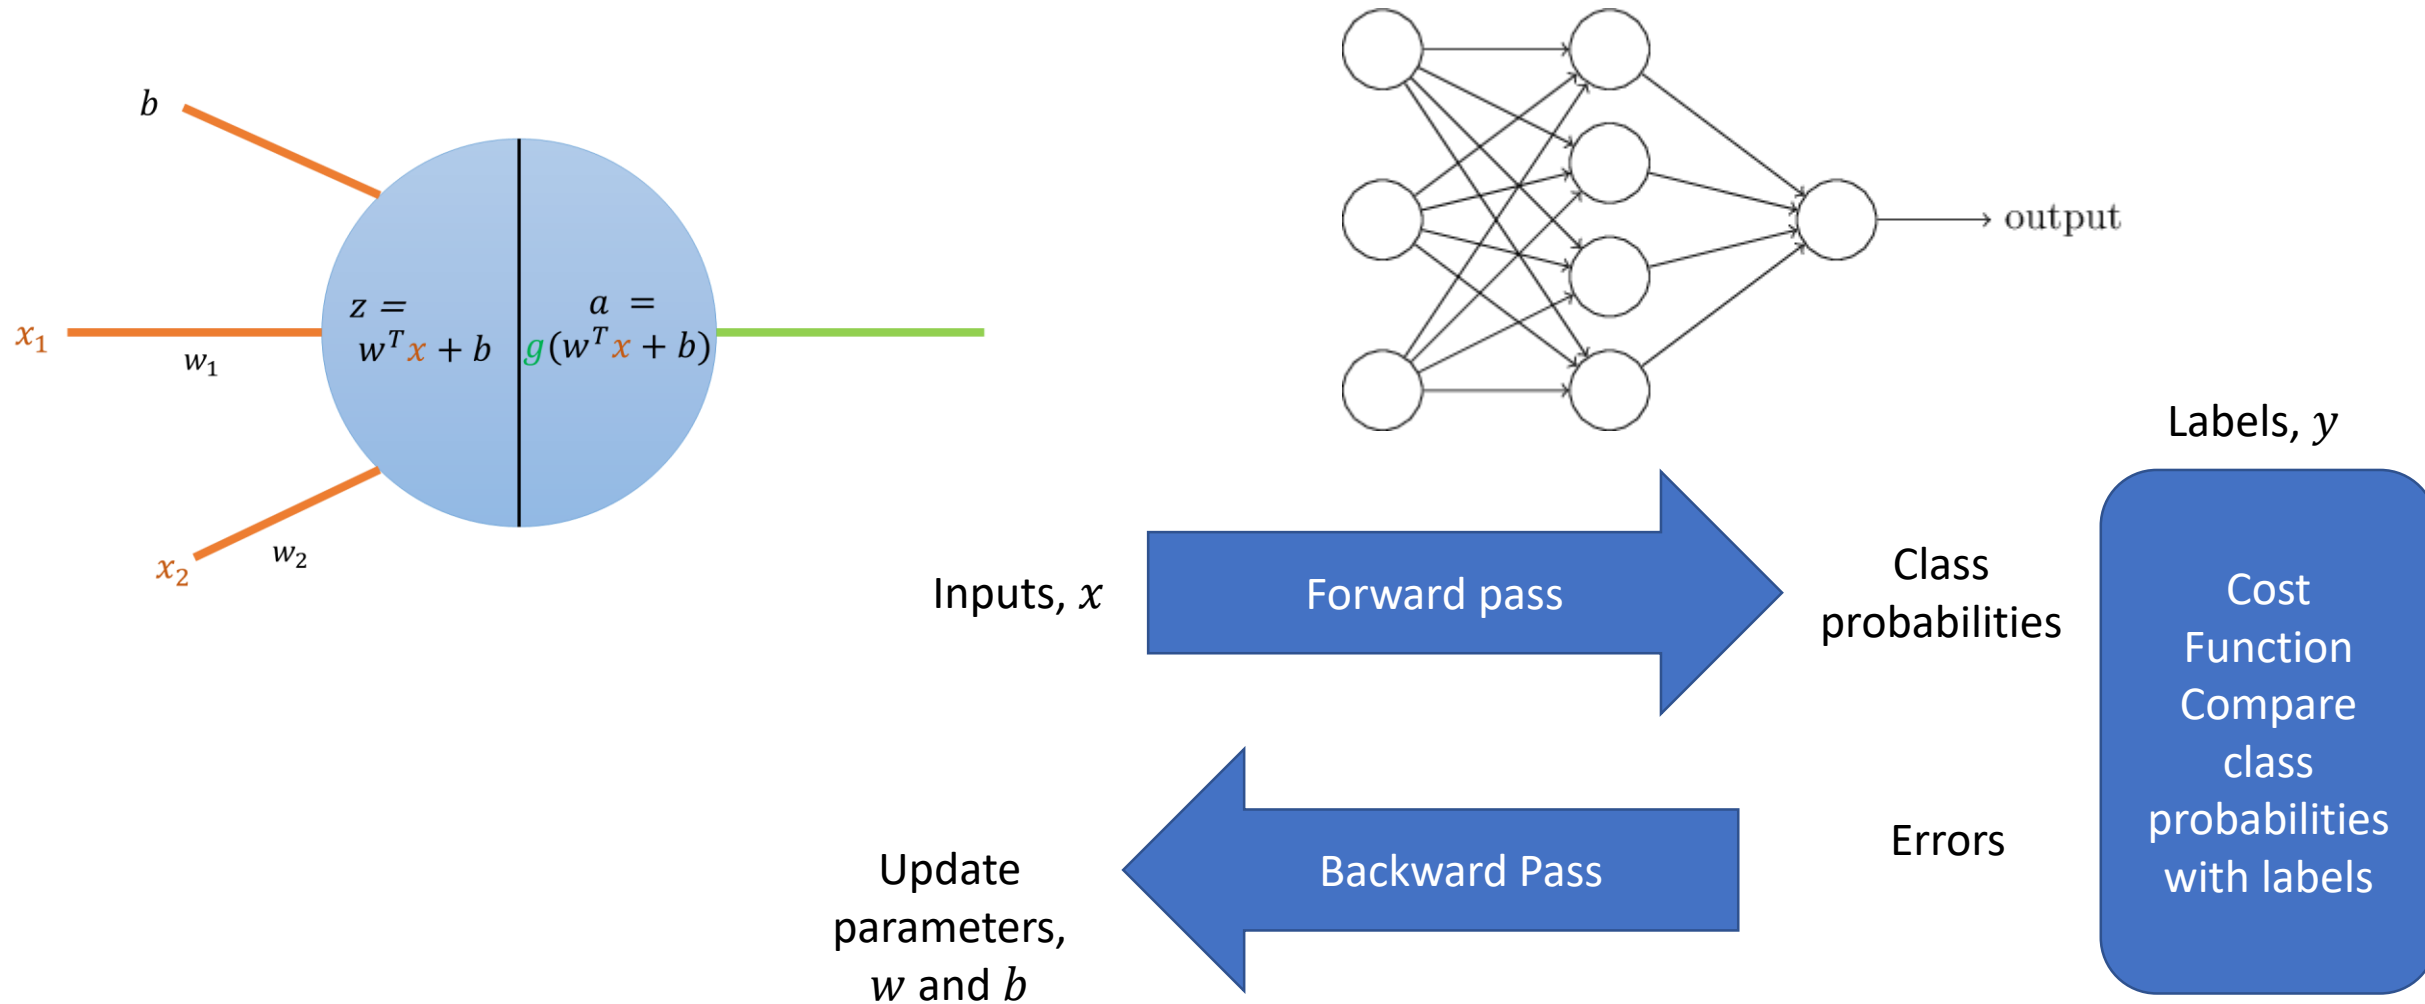
Logistic Regression to Neuron



Neural Network Examples

| Applications | Input | Output | Neural Network |
|----------------------|-------|--------|--------------------------|
| Housing Example | | | Simple NN |
| Iris Classification | | | Simple NN |
| Image classification | | | CNN |
| Tumor Segmentation | | | CNN |
| Speech translation | | | RNN / LSTM |
| Sentiment analysis | | | RNN / LSTM |
| Autonomous cars | | | Customized Architectures |

Neuron to Artificial neural network




Artificial Neural Networks

Input Data = \mathbf{x}, \mathbf{y}

Out Data = $\hat{\mathbf{y}}$


Fitness Function = $J(\mathbf{y}, \hat{\mathbf{y}})$

Weights and Biases = ???



Artificial Neural Networks

Forward Pass



$$Z^{[1]} = W^{[1]}X + b^{[1]}$$

$$A^{[1]} = \sigma(Z^{[1]})$$

$$Z^{[2]} = W^{[2]}A^{[1]} + b^{[2]}$$

$$A^{[2]} = \sigma(Z^{[2]})$$

$$X = \begin{bmatrix} | & | & | \\ x^{(1)} & x^{(2)} & x^{(m)} \\ | & | & | \end{bmatrix}$$


$$Z^{[1]} = \begin{bmatrix} | & | & | \\ z^{1} & z^{[1](2)} & z^{[1](m)} \\ | & | & | \end{bmatrix}$$

$$A^{[1]} = \begin{bmatrix} | & | & | \\ a^{1} & a^{[1](2)} & a^{[1](m)} \\ | & | & | \end{bmatrix}$$

```
Z1 = np.dot(W1, X) + b1
A1 = sigmoid(Z1)
Z2 = np.dot(W2, A1) + b2
A2 = sigmoid(Z2)
```

Artificial Neural Networks

Backpropagation



$$dZ^{[2]} = A^{[2]} - Y$$

$$dW^{[2]} = \frac{1}{m} dZ^{[2]} A^{[1]T}$$

$$db^{[2]} = \frac{1}{m} \sum dZ^{[2]}$$

$$dZ^{[1]} = W^{[2]T} dZ^{[2]} * \sigma'(Z^{[1]})$$



$$dW^{[1]} = dZ^{[1]} X^T$$

$$db^{[1]} = \frac{1}{m} \sum dZ^{[1]}$$

```
dZ2 = A2 - Y
dW2 = 1/m*(np.dot(dZ2,A1.T))
db2 = 1/m*(np.sum(dZ2, axis=1, keepdims=True))
dZ1 = np.dot(W2.T, dZ2)*(A1*(1-A1))
dW1 = 1/m*np.dot(dZ1, X.T)
db1 = (1/m) * np.sum(dZ1, axis=1, keepdims=True)
```

Artificial Neural Networks – Training



Classification Task



Initial random weights

Artificial Neural Networks – Training



Classification Task



Present a training instance / adjust the weights

Artificial Neural Networks – Training



Classification Task



Present a training instance / adjust the weights

Artificial Neural Networks – Training



Classification Task



Present a training instance / adjust the weights

Artificial Neural Networks – Training



Classification Task



Present a training instance / adjust the weights

Artificial Neural Networks – Training


Classification Task



Eventually


Artificial Neural Networks – Training and Testing

Train




| | | |
|-----|-----|-----|
| 73 | 73 | 101 |
| 101 | 134 | 101 |
| 134 | 73 | 134 |

Test




?


Artificial Neural Networks - Nonlinearity



The network will learn to classify points as belonging to one or the other.




ANN with 1 input layer and 1 output layer. The network simply tries to separate the two classes of data by dividing them with a line.




ANN with 1 hidden layer. It separates the data with a more complicated curve than a line.


Artificial Neural Networks – Why Go Deep?



1 hidden layer – 2 neurons



2 hidden layers – 4 neurons



3 hidden layers – 9 neurons

Background Math

- The Matrix Calculus You Need For Deep Learning

<https://arxiv.org/pdf/1802.01528.pdf>

Coding Assignment

- Logistic Regression
- Neural Network
 - Without using any libraries
 - Keras
 - TensorFlow
- Google co-lab

https://colab.research.google.com/drive/1H_F3Fgab3aTH1gBmXHTjlvarBheAJGvy